

Python Development Best Practices

Diogo Ribeiro
Data Scientist Lead
Portugal
<https://diogoribeiro7.github.io>
<https://github.com/DiogoRibeiro7>

November 13, 2025

Abstract

This book expands a practitioner-oriented article into a comprehensive handbook for intermediate and advanced Python developers. It explains why best practices matter, demonstrates how to apply them with idiomatic Python 3.12 examples, and describes how real teams use these techniques to ship reliable software. Each chapter blends theory, scenarios, code, and exercises so readers can internalise the material and adapt it to their own projects.

Contents

1	Introduction: Why Best Practices Matter	1
1.1	Chapter Overview	1
1.2	From Scripts to Systems	1
1.2.1	Scenario: The Overnight Script	2
1.3	Four Pillars of Professional Python	3
1.4	A Motivating Code Example	4
1.5	How to Use This Book	4
1.6	Summary	5
2	Code Style and Readability	7
2.1	Chapter Overview	7
2.2	PEP 8 as a Shared Vocabulary	7
2.3	Tooling Over Taste	8
2.4	Bad vs Good Example	8
2.5	Module-Sized Example	9
2.6	Scenario: Style Drift in a Release Crunch	9
2.7	Summary	10
3	Comments and Documentation	11
3.1	Chapter Overview	11
3.2	Docstrings that Educate	11
3.3	Inline Comments for Intent	12
3.4	Project Documentation Stack	12
3.5	Code Example: Documentation-Driven Development	13
3.6	Scenario: Docs Saved a Launch	13
3.7	Summary	13
4	Project Structure and Architecture	15
4.1	Chapter Overview	15
4.2	Repository Layout	15
4.3	Architectural Layers	16
4.4	Scenario: Refactoring a Monolith	16
4.5	Code Example: Service Layer Pattern	16
4.6	Summary	17

5 Dependencies and Environments	19
5.1 Chapter Overview	19
5.2 Virtual Environments Everywhere	19
5.3 Pinning and Locking	20
5.4 Example: <code>pyproject</code> Snippet	20
5.5 Scenario: Missing Dependency in Production	20
5.6 Summary	21
6 Type Hints and Static Analysis	23
6.1 Chapter Overview	23
6.2 Why Annotate	23
6.3 Protocols and Dataclasses	24
6.4 Static Analysis Workflow	24
6.5 Scenario: Type Hints Caught a Regression	25
6.6 Summary	25
7 Testing and Quality Assurance	27
7.1 Chapter Overview	27
7.2 Testing Pyramid	27
7.3 Pytest Example	28
7.4 Anti-Pattern: Untestable Code	28
7.5 Scenario: Catching a Production Bug	29
7.6 Summary	29
8 Error Handling and Logging	31
8.1 Chapter Overview	31
8.2 Domain-Specific Exceptions	31
8.3 Structured Logging	32
8.4 Bad vs Good Error Handling	32
8.5 Scenario: Incident Response	33
8.6 Summary	33
9 Configuration and Secrets Management	35
9.1 Chapter Overview	35
9.2 Typed Configuration Loader	35
9.3 Secrets Discipline	36
9.4 Scenario: Leaked Sandbox Key	36
9.5 Summary	36
10 Performance and Scalability	39
10.1 Chapter Overview	39
10.2 Profiling Before Optimising	39
10.3 Data Structures	39
10.4 Concurrency Choices	40
10.5 Scenario: Scaling a Report Generator	40
10.6 Summary	41

11 Security for Python Developers	43
11.1 Chapter Overview	43
11.2 Dependency Hygiene	43
11.3 Input Validation	43
11.4 Safe File Handling	44
11.5 Scenario: Dependency Supply-Chain Attack	44
11.6 Summary	45
12 Packaging, Distribution, and Versioning	47
12.1 Chapter Overview	47
12.2 Build Real Packages	47
12.3 Entry Points	48
12.4 Versioning Discipline	48
12.5 Scenario: Coordinated Release Train	48
12.6 Summary	49
13 Collaboration, Git, and CI/CD	51
13.1 Chapter Overview	51
13.2 Git Hygiene	51
13.3 Code Review Culture	51
13.4 Automation and Hooks	52
13.5 Scenario: CI as Gatekeeper	52
13.6 Summary	53
14 Case Studies and Patterns in Real Projects	55
14.1 Chapter Overview	55
14.2 Billing Platform Modernisation	55
14.3 Data Pipeline Hardening	56
14.4 Future Work	56
14.5 Summary	56
A Checklists, Templates, and Further Reading	59
A.1 Chapter Overview	59
A.2 Launch Checklist	59
A.3 Template Repository	59
A.4 Further Reading	60
A.5 Summary	60

Chapter 1

Introduction: Why Best Practices Matter

1.1 Chapter Overview

Professional Python development rarely fails because of missing syntax knowledge. Most problems emerge when the surrounding practices—testing, structure, automation, culture—are weak or inconsistent. This opening chapter reframes best practices as strategic tools for reducing risk and enabling collaboration.

```
from __future__ import annotations

from dataclasses import dataclass


@dataclass(slots=True)
class PracticeScore:
    """Represent how consistently a team applies a given practice."""

    name: str
    health: float  # 0-1 scale
    blocking_incidents: int

    def needs_exec_attention(score: PracticeScore) -> bool:
        """Escalate when weakened practices repeatedly cause incidents."""
        return score.health < 0.6 and score.blocking_incidents >= 2
```

Listing 1.1: Tracking practice adoption across a portfolio

Teams often add similar lightweight metrics to post-incident reviews so they can link failures back to neglected practices instead of blaming individuals.

1.2 From Scripts to Systems

Many engineers fall in love with Python by automating a tedious task. The moment that script becomes a service, library, or data pipeline, the definition of success changes. Colleagues need repeatable environments, reliable releases, and discoverable behaviour. Ignoring those expectations leads to outages that are harder to diagnose than the original problem.

```

from __future__ import annotations

from dataclasses import dataclass
from pathlib import Path


@dataclass(slots=True)
class SyncJobConfig:
    """Describe the inputs required to run the nightly sync job."""

    input_dir: Path
    output_dir: Path
    dry_run: bool = False

    def run_sync_job(config: SyncJobConfig) -> int:
        """Convert ad-hoc filesystem work into a composable, testable unit."""
        processed = 0
        for source in config.input_dir.glob("*.csv"):
            destination = config.output_dir / source.name
            if not config.dry_run:
                destination.write_text(source.read_text(encoding="utf-8"),
                                      encoding="utf-8")
            processed += 1
        return processed

```

Listing 1.2: Promoting an exploratory script into a service-friendly module

Bad scripts hide configuration inside global constants; production-ready modules model dependencies explicitly the way ‘SyncJobConfig’ does, making it obvious how to parameterise jobs for CI or data replay.

1.2.1 Scenario: The Overnight Script

A support engineer wrote a quick file-parsing script that was then scheduled to run nightly on production data. Weeks later an upstream schema change broke the script silently and the company did not notice until invoices failed. The fix involved adding tests, documentation, and alerts—essentially retrofitting everything a well-structured project would have had from day one. The lesson: treat every piece of code as a future dependency.

```

def nightly_parse(path: str) -> list[str]:
    """Process rows but fail silently when columns change."""
    results = []
    with open(path, encoding="utf-8") as handle:
        for line in handle:
            parts = line.strip().split(",")
            results.append(parts[3]) # hard-coded column index
    return results

```

Listing 1.3: Anti-pattern: implicit schema assumptions

```

from __future__ import annotations

from pathlib import Path

```

```

EXPECTED_HEADERS = ("invoice_id", "customer_id", "status", "total")

def parse_invoices(path: Path) -> list[dict[str, str]]:
    """Validate incoming files so incidents surface immediately."""
    lines = path.read_text(encoding="utf-8").splitlines()
    header = tuple(lines[0].split(","))
    if header != EXPECTED_HEADERS:
        raise ValueError(f"Unexpected schema: {header}")
    return [
        dict(zip(EXPECTED_HEADERS, row.split(","), strict=True))
        for row in lines[1:]
    ]

```

Listing 1.4: Resilient parser with validation hooks

When the team deployed the validated parser, they also routed failures to PagerDuty, transforming a silent data loss issue into a ten-minute debugging session.

1.3 Four Pillars of Professional Python

Best practices reinforce four outcomes:

Readability Code tells a clear story so reviewers and incident responders can reason quickly.

Reliability Tests, type hints, and logging ensure behaviour is predictable and regressions are caught early.

Maintainability Modular design, versioning, and automation keep change costs low even as teams grow.

Reproducibility Environments, data, and configuration can be recreated on fresh machines, enabling confident deployments.

```

from __future__ import annotations

from enum import Enum, auto

class Pillar(Enum):
    READABILITY = auto()
    RELIABILITY = auto()
    MAINTAINABILITY = auto()
    REPRODUCIBILITY = auto()

def missing_pillars(observed: set[Pillar]) -> list[Pillar]:
    """Return the set of pillars not satisfied by a service review."""
    return [pillar for pillar in Pillar if pillar not in observed]

```

Listing 1.5: Capturing pillar health as structured data

Teams often embed functions like `missing_pillars` inside CI linters so pull requests fail when key artefacts—tests, docs, deployment manifests—are absent.

1.4 A Motivating Code Example

Listing 1.6 juxtaposes an exploratory script with a production-ready function. The logic is the same, but the maintainable version adds validation, typing, and logging hooks to support future owners.

```
from __future__ import annotations

from collections.abc import Iterable

def total_active_users(user_counts: Iterable[int]) -> int:
    """Return the sum of daily active users after validating inputs."""
    totals = []
    for count in user_counts:
        if count < 0:
            raise ValueError("daily active users cannot be negative")
        totals.append(count)
    return sum(totals)
```

Listing 1.6: Transforming a script into a reusable function

1.5 How to Use This Book

Read sequentially if you are building a new codebase. Dip into specific chapters when improving an existing system. Each chapter ends with exercises to reinforce the concepts and to provide prompts for team discussions.

```
from __future__ import annotations

from dataclasses import dataclass


@dataclass(slots=True)
class Chapter:
    """Metadata describing how a team might prioritise chapters."""

    name: str
    focus: str
    effort_hours: int

    def pick_next_chapter(chapters: list[Chapter], topic: str) -> Chapter:
        """Select the shortest chapter that matches the requested topic."""
        candidates = [chapter for chapter in chapters if chapter.focus == topic]
        return min(candidates, key=lambda chapter: chapter.effort_hours)
```

Listing 1.7: Generating a study plan from chapter metadata

Embedding chapter planning in onboarding tooling helps managers pair new hires with the sections most relevant to current incidents or projects.

1.6 Summary

Best practices succeed when teams treat every script as a future dependency, pursue the four pillars consistently, and lean on automation plus documentation to scale collaboration. The rest of the book dives into each pillar with practical guidance.

```
from __future__ import annotations

from collections.abc import Iterable


def summarize_retrospective(findings: Iterable[str]) -> dict[str, list[str]]:
    """Group retrospective findings by pillar for postmortem dashboards."""
    summary: dict[str, list[str]] = {"people": [], "process": [], "tooling": []}
    for finding in findings:
        bucket = "tooling" if "automation" in finding else "process"
        if "on-call" in finding:
            bucket = "people"
        summary[bucket].append(finding)
    return summary
```

Listing 1.8: Turning retrospectives into action items

Exercises

1. Identify a script in your organisation that quietly became a service. List the missing practices (tests, logging, configuration) and sketch a plan to add them.
2. Rewrite a quick-and-dirty function by applying the patterns from Listing 1.6. What impeded readability?
3. Interview a teammate about a recent incident. Which of the four pillars would have prevented or shortened it?
4. Draft a one-page document explaining how new hires should navigate your repository. Share it with the team and capture feedback.
5. Build a checklist for scripts that might graduate to production and store it in version control.

Chapter 2

Code Style and Readability

2.1 Chapter Overview

Style is not mere aesthetics; it is a communication contract. When code follows shared conventions, teams spend less time deciphering formatting and more time reasoning about behaviour.

```
from __future__ import annotations

from dataclasses import dataclass


@dataclass(slots=True)
class StyleScore:
    """Lightweight structure for tracking readability drift."""

    module: str
    naming_rules_passed: bool
    formatter_ran: bool
```

Listing 2.1: Capturing style rule adherence

2.2 PEP 8 as a Shared Vocabulary

PEP 8 codifies naming, indentation, and whitespace rules. Adhering to it means any Python developer can jump into the repository without a style orientation session. When exceptions are necessary—long SQL strings or generated code—document the rationale in the README or a style guide.

```
from __future__ import annotations


def enforce_snake_case(names: list[str]) -> list[str]:
    """Return identifiers that require refactoring."""
    return [name for name in names if any(ch.isupper() for ch in name)]
```

Listing 2.2: Encoding naming rules in review helpers

Many teams run helpers like `enforce_snake_case` in review bots so regression is impossible without an explicit waiver.

2.3 Tooling Over Taste

Automate style enforcement to reduce cognitive load. `black` or `ruff format` handle formatting, `ruff` enforces linting rules, and `isort` (or `ruff`'s import subsystem) maintains deterministic import order. Configure the tools in `pyproject.toml` and run them via `pre-commit` so every change goes through the same gate.

```
from __future__ import annotations

import subprocess


def run_formatter() -> None:
    """Exit the build early when the formatter changes files."""
    result = subprocess.run(["ruff", "format", "--check", "."], check=False)
    if result.returncode != 0:
        msg = "Formatting drift detected; run 'ruff format' locally."
        raise SystemExit(msg)
```

Listing 2.3: Guarding formatters inside CI scripts

2.4 Bad vs Good Example

Listing 2.4 and Listing 2.5 show how naming and structure transform comprehension. Both compute a discounted total, but the improved version reveals intent immediately.

```
def calc(d, p):
    t = 0
    for x in d:
        t += x
    if p:
        t -= t * p
    return t
```

Listing 2.4: Anti-pattern: inconsistent naming and hidden branching

```
from __future__ import annotations


def calculate_discounted_total(prices: list[float], discount: float) -> float:
    """Return the discounted total after validating all inputs."""
    if discount < 0 or discount > 1:
        raise ValueError("discount must be a percentage between 0 and 1")

    if any(price < 0 for price in prices):
        raise ValueError("prices must be non-negative")

    subtotal = sum(prices)
    return subtotal * (1 - discount)
```

Listing 2.5: Readable implementation with validation

2.5 Module-Sized Example

Listing 2.6 demonstrates how naming, docstrings, and dataclasses combine in a realistic module that a billing service could ship.

```
from __future__ import annotations

from dataclasses import dataclass
from datetime import datetime
from decimal import Decimal, ROUND_HALF_UP
from typing import Iterable

@dataclass(frozen=True)
class InvoiceLine:
    description: str
    quantity: int
    unit_price: Decimal

    def total(self) -> Decimal:
        """Return a rounded line total."""
        subtotal = self.unit_price * self.quantity
        return subtotal.quantize(Decimal("0.01"), rounding=ROUND_HALF_UP)

def summarise_invoice(lines: Iterable[InvoiceLine], issued_at: datetime) \
-> dict[str, object]:
    """Produce a payload for downstream email, PDF, or API consumers."""
    totals = [line.total() for line in lines]
    return {
        "issued_at": issued_at.isoformat(),
        "line_count": len(totals),
        "subtotal": sum(totals),
        "currency": "USD",
    }
```

Listing 2.6: Invoice summariser for a billing engine

2.6 Scenario: Style Drift in a Release Crunch

During a tight deadline, a payments team disabled the formatter because it produced merge conflicts with an experimental branch. Within two weeks, code reviews were bogged down by spacing debates and subtle bugs introduced by inconsistent imports. After the release, they reinstated `pre-commit` hooks and added `ruff` checks to CI to prevent future drift.

```
from __future__ import annotations

from pathlib import Path

def detect_style_drift(diff_root: Path) -> list[Path]:
    """Spot modules that violate formatting during emergency fixes."""
    offenders = []
```

```

for path in diff_root.rglob("*.py"):
    if "\t" in path.read_text(encoding="utf-8"):
        offenders.append(path)
return offenders

```

Listing 2.7: Detecting inconsistent formatting in a pull request

When the crunch was over, the team ran `detect_style_drift` against every hotfix branch, opened follow-up issues for each offender, and cleared the debt before the next release.

2.7 Summary

Consistent style grows from agreed rules, automated enforcement, and shared examples that clarify expectations. Keep the tooling always-on so humans focus on logic rather than indentation debates.

```

from __future__ import annotations

def summarize_violations(violations: dict[str, int]) -> str:
    """Return a short sentence for incident reviews."""
    worst_offender = max(violations, key=violations.get)
    total = sum(violations.values())
    return f"{total} style issues detected; {worst_offender} needs a
           cleanup."

```

Listing 2.8: Summarising style violations for dashboards

Exercises

1. Format a messy module manually, then re-run `black`. Note every change the formatter made and decide which rules you would not have caught.
2. Using Listing 2.4, annotate each readability issue and propose a rule that would catch it automatically.
3. Configure `ruff` to enforce import sorting. Break the rule intentionally and observe the linter output.
4. Pair with a teammate and review a 200-line pull request. Track how often naming or formatting slowed understanding. Compare notes and create a shared checklist.
5. Extend your team's style guide with two domain-specific conventions (e.g., how to name dataclass factories) and circulate the update.
6. Refactor the snippet below into readable code with docstrings and typing.

```

def calc(x, y):
    return sum([a * y for a in x if a > 0])

```

Listing 2.9: Exercise: refactor for clarity

7. Evaluate a formatter configuration that lives in `pyproject.toml`. Explain which defaults you overrode and why they matter to your domain.

Chapter 3

Comments and Documentation

3.1 Chapter Overview

Well-structured code minimises the need for comments, yet documentation remains essential for sharing context, intent, and processes.

```
from __future__ import annotations

from dataclasses import dataclass
from datetime import datetime

@dataclass(slots=True)
class DocumentStatus:
    """Keep tabs on which guides need attention."""

    path: str
    last_reviewed: datetime

    def needs_refresh(status: DocumentStatus, *, months: int = 6) -> bool:
        """Flag documents that have gone stale."""
        delta = datetime.now(tz=status.last_reviewed.tzinfo) - status.
            last_reviewed
        return delta.days > months * 30
```

Listing 3.1: Tracking documentation freshness

3.2 Docstrings that Educate

Docstrings should answer three questions: what does this item do, what are its inputs and outputs, and when does it raise exceptions? Follow the "one-line summary + details" format so tools like Sphinx can render the text neatly.

```
from __future__ import annotations

def parse_currency(value: str) -> int:
```

```
"""Return cents for a currency string; raise ValueError on malformed
input."""
if not value.startswith("$"):
    raise ValueError("Currency must start with a dollar sign.")
dollars, cents = value[1:].split(".")
return int(dollars) * 100 + int(cents)
```

Listing 3.2: Docstring capturing behaviour and failure modes

3.3 Inline Comments for Intent

Inline comments should explain *why*, not *what*. They are ideal for noting workarounds, referencing bug tickets, or clarifying domain constraints. Delete them once they become obsolete to avoid misleading readers.

```
from __future__ import annotations

from datetime import UTC, datetime

def next_billing_window(now: datetime) -> datetime:
    """Return the start of the next billing window."""
    window = now.astimezone(UTC).replace(minute=0, second=0, microsecond=0)
    # Stripe batches reconcile at 5-minute intervals, so align to that
    # cadence.
    return window.replace(minute=(window.minute // 5 + 1) * 5)
```

Listing 3.3: Inline comments that explain surprising workarounds

3.4 Project Documentation Stack

Combine lightweight Markdown guides with generated API references. Example structure: `README.md` for onboarding, `docs/architecture.md` for high-level design, Sphinx or MkDocs for API references, and runbooks for deployment and incident response.

```
from __future__ import annotations

from pathlib import Path

def ensure_docs_index(root: Path) -> None:
    """Create a landing page that references every Markdown guide."""
    guides = sorted(root.glob("*.md"))
    body = "\n".join(f"- [{guide.stem}]({guide.name})" for guide in guides)
    (root / "index.md").write_text(body, encoding="utf-8")
```

Listing 3.4: Generating a docs index programmatically

3.5 Code Example: Documentation-Driven Development

Listing 3.5 demonstrates an API that uses docstrings and inline comments to highlight non-obvious behaviour.

```
from __future__ import annotations

from datetime import UTC, datetime

def anonymise_timestamp(raw_timestamp: str) -> datetime:
    """Return an hourly bucket so analytics cannot deanonymise individual
    events."""
    timestamp = datetime.fromisoformat(raw_timestamp).astimezone(UTC)
    # We discard minutes/seconds to comply with the analytics retention
    # policy.
    return timestamp.replace(minute=0, second=0, microsecond=0)
```

Listing 3.5: Docstring and inline comment illustrating intent

3.6 Scenario: Docs Saved a Launch

A healthcare startup had to demonstrate auditability to regulators. Because the team maintained up-to-date architecture diagrams and onboarding walkthroughs, they produced the requested evidence in hours instead of weeks. The documentation discipline paid off even though the company had never planned for a government review.

```
from __future__ import annotations

from pathlib import Path

def gather_audit_artifacts(root: Path) -> list[Path]:
    """Return runbooks and diagrams needed during regulated reviews."""
    return [path for path in root.rglob("*") if path.suffix in {".md", ".drawio"}]
```

Listing 3.6: Collecting evidence links quickly

3.7 Summary

Documentation succeeds when it explains intent, evolves with the system, and exists at multiple layers—from docstrings to architecture guides and runbooks.

```
from __future__ import annotations

def summarize_gaps(gaps: dict[str, list[str]]) -> str:
    """Turn gap analysis into a concise statement for leadership."""
    missing = ", ".join(sorted(gaps))
    return f"Docs need updates for {missing}; owners assigned."
```

Listing 3.7: Summarising documentation gaps

Exercises

1. Choose an undocumented function and write a docstring that follows the style in Listing 3.5.
2. Create an inline comment explaining a surprising line of code, then refactor until the comment is unnecessary.
3. Generate API documentation with Sphinx for a small module. What metadata is missing?
4. Draft a runbook for a scheduled job. Include inputs, expected outputs, and failure procedures.
5. Interview a new hire about confusing documentation gaps and log the findings as GitHub issues.
6. Use the snippet below to explain how you would document the edge cases.

```
def compute(offset: int) -> int:
    if offset < 0:
        return 0
    return offset * 2
```

Listing 3.8: Exercise: document behaviour

7. Sketch a documentation stack diagram that shows how README files, architecture records, and runbooks link to each other. Identify who updates each artefact.

Chapter 4

Project Structure and Architecture

4.1 Chapter Overview

A thoughtful structure reinforces boundaries and accelerates development. This chapter covers folder layout, layering, and refactoring strategies.

```
from __future__ import annotations

from dataclasses import dataclass

@dataclass(slots=True)
class ModuleBoundary:
    """Describe how code is grouped and who maintains it."""

    path: str
    owner_team: str
    depends_on: list[str]
```

Listing 4.1: Capturing ownership of directories

4.2 Repository Layout

Place application code in `src/`, tests in `tests/`, documentation in `docs/`, and experimentation in `experiments/`. The goal is to make every directory self-explanatory.

```
python-best-practices/
|-- pyproject.toml
|-- README.md
|-- src/
|   '-- billing/
|       |-- __init__.py
|       |-- invoices.py
|       '-- repositories.py
|-- tests/
|   '-- test_invoices.py
`-- docs/
    '-- architecture.md
```

Listing 4.2: Predictable repository layout

4.3 Architectural Layers

Layered architectures separate domain logic from infrastructure. A typical service contains domain models, repositories, adapters, and entry points (CLI, HTTP, messaging). Each layer depends inward, making it easier to test and swap components.

```
from __future__ import annotations

from dataclasses import dataclass


@dataclass(slots=True)
class Command:
    """DTO that carries user intent from the presentation layer."""

    user_id: str
    action: str
```

Listing 4.3: Defining explicit boundaries between layers

4.4 Scenario: Refactoring a Monolith

A logistics company inherited a single-file monolith that processed shipments, handled notifications, and exposed a CLI. By splitting responsibilities into modules that matched their architecture diagram, they unlocked parallel development. Teams could now work on inventory, billing, or reporting without stepping on each other's toes.

```
def process_all() -> None:
    """Monolithic script that mixes unrelated behaviours."""
    load_inventory()
    notify_customers()
    run_cli()
```

Listing 4.4: Coarse-grained monolith vs refactored entry point

```
from __future__ import annotations


def orchestrate() -> None:
    """Dispatch into modules that mirror architecture boundaries."""
    inventory.run_pipeline()
    notifications.send_digest()
    cli.run()
```

Listing 4.5: Modular orchestration after refactor

4.5 Code Example: Service Layer Pattern

Listing 4.6 shows a service that decouples use cases from persistence.

```
from __future__ import annotations
```

```

from dataclasses import dataclass
from typing import Protocol
from uuid import UUID

class UserGateway(Protocol):
    def deactivate(self, user_id: UUID) -> None: ...

class UserNotFoundError(RuntimeError):
    """Raised when a user identifier does not exist."""

@dataclass(slots=True)
class UserService:
    gateway: UserGateway

    def deactivate_user(self, user_id: UUID) -> None:
        try:
            self.gateway.deactivate(user_id)
        except UserNotFoundError as exc:
            raise ValueError(f"Unknown user {user_id}") from exc

```

Listing 4.6: Service layer separating domain and persistence

4.6 Summary

Let the directory tree mirror the architecture, keep responsibilities separated by layers, and refactor incrementally so structure evolves alongside the product.

```

from __future__ import annotations

def assert_dependency_rule(module: str, allowed: set[str], imports: set[str]) -> None:
    """Guardrails to ensure layers only depend inward."""
    illegal = imports - allowed
    if illegal:
        raise AssertionError(f"{module} cannot import {illegal}")

```

Listing 4.7: Asserting architectural contracts in tests

Exercises

1. Draw an architecture diagram of your current project. How well does the folder structure reflect that diagram?
2. Apply the service-layer pattern from Listing 4.6 to another use case (e.g., invoice approval).
3. Identify a module that performs both domain and persistence work. Split it into two files and note how tests change.

4. Design a strategy for handling cross-cutting concerns (validation, logging) without violating layering.
5. Propose a directory naming convention for experiments or spikes and document it.
6. Use the pseudo-monolith snippet in Listing 4.4 as input and sketch a refactoring plan that introduces adapters, services, and repositories.

```
from __future__ import annotations

from pathlib import Path


def validate_layout(root: Path) -> None:
    """Ensure expected top-level directories exist."""
    required = {"src", "tests", "docs"}
    missing = required - {path.name for path in root.iterdir() if path.
        is_dir()}
    if missing:
        raise SystemExit(f"Missing required directories: {sorted(missing)}")
```

Listing 4.8: Validating repository layout in CI

Chapter 5

Dependencies and Environments

5.1 Chapter Overview

Environment drift is a leading cause of “works on my machine” bugs. This chapter explains how to manage dependencies systematically.

```
from __future__ import annotations

from dataclasses import dataclass

@dataclass(slots=True)
class EnvironmentReport:
    """Track interpreter, dependency lock, and platform data."""

    python_version: str
    lockfile_hash: str
    platform: str
```

Listing 5.1: Capturing environment information

5.2 Virtual Environments Everywhere

Use `python -m venv`, `uv`, or `pipenv` to isolate dependencies per project. Document activation commands and add helper scripts that install tooling in one step.

```
from __future__ import annotations

import subprocess
from pathlib import Path


def create_environment(venv_dir: Path) -> None:
    """Create and populate a deterministic virtual environment."""
    subprocess.run(["python", "-m", "venv", str(venv_dir)], check=True)
    subprocess.run([str(venv_dir / "Scripts" / "python"), "-m", "pip", "install", "-U", "pip"], check=True)
```

Listing 5.2: Bootstrapping a virtual environment programmatically

5.3 Pinning and Locking

Declare dependencies in `pyproject.toml` and generate lock files with `uv lock` or `pip-compile`. Commit the lock files so CI and production match local development.

```
from __future__ import annotations

import hashlib
from pathlib import Path


def lock_hash(path: Path) -> str:
    """Return the hash used to verify reproducible installs."""
    return hashlib.sha256(path.read_bytes()).hexdigest()
```

Listing 5.3: Verifying lock files before deployment

5.4 Example: `pyproject` Snippet

```
[project]
name = "billing-service"
version = "0.4.0"
requires-python = ">=3.12"
dependencies = [
    "pydantic>=2.5",
    "ruff==0.6.8",
]

[project.optional-dependencies]
dev = ["pytest>=8.3", "mypy>=1.11", "hypothesis>=6.99"]
```

Listing 5.4: `pyproject` snippet with optional dependencies

5.5 Scenario: Missing Dependency in Production

An e-commerce team relied on a globally installed CLI utility that was present on their laptops but not on CI servers. A deployment failed hours before Black Friday. The root cause was a missing dependency specification. After the incident, they moved all tooling into `pyproject.toml` and added smoke tests to CI that bootstrap environments from scratch.

```
from __future__ import annotations

from importlib import metadata


def assert_tool_declared(package: str) -> None:
    """Fail fast when required tooling is absent."""
    try:
        metadata.version(package)
    except metadata.PackageNotFoundError as exc:
        raise RuntimeError(f"Missing dependency: {package}") from exc
```

Listing 5.5: Ensuring tools are declared explicitly

5.6 Summary

Reliable deployments depend on isolated environments, locked dependency graphs, and CI pipelines that recreate installations exactly as production does.

```
from __future__ import annotations

def summarize_dependencies(outdated: list[str]) -> str:
    """Provide leadership-friendly summary of update backlog."""
    if not outdated:
        return "All dependencies are current."
    return f"{len(outdated)} packages require upgrades: {', '.join(
        outdated[:5])}..."
```

Listing 5.6: Summarising dependency health for dashboards

Exercises

1. Regenerate your lock file using `uv` or `pip-compile`. How many transitive dependencies changed?
2. Write a shell script that recreates the project environment from scratch. Run it on a clean machine or container.
3. Audit your repository for globally installed tools. Add them to `pyproject` extras.
4. Introduce a deliberate dependency conflict and run `pip check` to observe the failure.
5. Design a process for handling emergency dependency upgrades triggered by security advisories.
6. Extend Listing 5.3 to compare hash values between CI and production deployments. Explain how you would alert on mismatches.

```
from __future__ import annotations

import subprocess

def install_group(group: str) -> None:
    """Install an optional dependency group from pyproject metadata."""
    subprocess.run(["uv", "pip", "install", f".[{group}]"], check=True)
```

Listing 5.7: Install dependencies from the defined groups

Chapter 6

Type Hints and Static Analysis

6.1 Chapter Overview

Type hints capture intent and enable automated reasoning. Combined with static analyzers, they prevent entire classes of bugs.

```
from __future__ import annotations

def annotation_coverage(total_objects: int, typed_objects: int) -> float:
    """Return the proportion of typed callables and modules."""
    return round(typed_objects / total_objects, 2)
```

Listing 6.1: Capturing annotation coverage

6.2 Why Annotate

Annotations clarify contracts for both humans and tools. They document mutability, optional values, and expected shapes. Their primary benefit is catching mismatches at development time rather than in production.

```
from __future__ import annotations

from dataclasses import dataclass
from datetime import datetime

@dataclass(slots=True)
class InvoiceDraft:
    """Represent an invoice that may or may not be finalised."""

    issued_at: datetime | None
    total: int

    def finalize(draft: InvoiceDraft, *, issued_at: datetime) -> InvoiceDraft:
        """Produce a finalized invoice with explicit timestamps."""
        return InvoiceDraft(issued_at=issued_at, total=draft.total)
```

Listing 6.2: Capturing optional state via annotations

6.3 Protocols and Dataclasses

Protocols describe structural typing, allowing you to express behaviour without inheritance. Dataclasses handle boilerplate while keeping types explicit.

```
from __future__ import annotations

from collections.abc import Iterable
from dataclasses import dataclass
from datetime import datetime, timedelta
from typing import Protocol


class Schedulable(Protocol):
    priority: int

    def execute(self) -> None: ...


@dataclass(slots=True)
class ScheduledJob:
    job: Schedulable
    eta: datetime


def schedule_jobs(jobs: Iterable[Schedulable], *, start: datetime) -> list[ScheduledJob]:
    """Order jobs by priority and assign execution windows."""
    sorted_jobs = sorted(jobs, key=lambda job: job.priority, reverse=True)
    schedule: list[ScheduledJob] = []
    for offset, job in enumerate(sorted_jobs):
        schedule.append(ScheduledJob(job=job, eta=start + timedelta(
            minutes=offset)))
    return schedule
```

Listing 6.3: Scheduling interface using protocols

6.4 Static Analysis Workflow

Run `mypy` or `pyright` in CI alongside `ruff`. Treat warnings as actionable feedback. If a suppression is necessary, add a comment explaining why.

```
from __future__ import annotations

import subprocess


def run_type_checks() -> None:
    """Invoke mypy with the strict profile used in CI."""
    subprocess.run(["mypy", "--strict", "src"], check=True)
```

Listing 6.4: Running `mypy` with strict settings programmatically

6.5 Scenario: Type Hints Caught a Regression

A data platform refactored a parser to support new file formats. Type hints and `mypy` flagged that the new code returned `None` in a branch where downstream code expected a `str`. The regression never reached staging.

```
from __future__ import annotations

def parse_status(raw: str) -> str:
    """Return a validated status string."""
    match raw.lower():
        case "paid" | "pending":
            return raw.lower()
        case _:
            raise ValueError(f"Unknown status: {raw}")
```

Listing 6.5: Regression prevented by precise return types

6.6 Summary

Treat annotations as executable documentation, layer protocols and dataclasses to express behaviour, and rely on static analyzers for immediate feedback during reviews and CI.

```
from __future__ import annotations

def summarize_type_errors(errors: list[str]) -> str:
    """Condense mypy output for leadership updates."""
    if not errors:
        return "Type checks clean; no regressions."
    return f"{len(errors)} type errors remain; triage blockers first."
```

Listing 6.6: Summarising type-check results

Exercises

1. Add type hints to an untyped module and run `mypy`. Document each warning and the fix.
2. Convert a nominal interface into a `Protocol` and note how it simplifies testing.
3. Configure `pyright` or `mypy` in `pre-commit`. Measure the runtime impact.
4. Identify a place where you rely on `Any`. Replace it with a precise type.
5. Debate with your team when to use `TypedDict`, `dataclass`, or `attrs`.

Chapter 7

Testing and Quality Assurance

7.1 Chapter Overview

Testing gives confidence that code behaves as designed. This chapter covers strategy, tooling, and common pitfalls.

```
from __future__ import annotations

def test_success_rate(passed: int, total: int) -> float:
    """Return percentage of passing tests across the suite."""
    return round((passed / total) * 100, 2)
```

Listing 7.1: Capturing test suite health

7.2 Testing Pyramid

Think of tests as layers:

- Unit tests: fast, pure logic.
- Integration tests: real dependencies (databases, APIs).
- End-to-end tests: simulate user workflows.

Balance investment so feedback stays fast but coverage remains meaningful.

```
from __future__ import annotations

def classify_test(duration_seconds: float) -> str:
    """Return the pyramid layer a test belongs to."""
    if duration_seconds < 0.2:
        return "unit"
    if duration_seconds < 5:
        return "integration"
    return "end_to_end"
```

Listing 7.2: Defining representative tests per layer

7.3 Pytest Example

```

from datetime import UTC, datetime
from decimal import Decimal

import pytest

from billing.invoices import InvoiceLine, summarise_invoice


@pytest.fixture
def sample_lines() -> list[InvoiceLine]:
    return [
        InvoiceLine(description="Subscription", quantity=1, unit_price=Decimal("49.00")),
        InvoiceLine(description="Support", quantity=3, unit_price=Decimal("15.00")),
    ]


def test_summarise_invoice_counts_lines(sample_lines, freezer):
    freezer.move_to("2024-05-01T12:00:00Z")
    payload = summarise_invoice(sample_lines, issued_at=datetime.now(tz=UTC))
    assert payload["line_count"] == 2
    assert payload["subtotal"] == Decimal("94.00")

```

Listing 7.3: Pytest module with fixtures and time control

7.4 Anti-Pattern: Untestable Code

```

import os
import shutil
from uuid import uuid4

import boto3


def sync_everything() -> None:
    client = boto3.client("s3")
    for path in os.listdir("/tmp"):
        if path.endswith(".json"):
            client.upload_file(path, "analytics-bucket", f"imports/{uuid4()}")
            shutil.move(path, "/archive")

```

Listing 7.4: Deeply coupled function that defies testing

Break the function into collaborators that can be faked in tests. Inject the S3 client and filesystem wrapper rather than instantiating them inline.

7.5 Scenario: Catching a Production Bug

A fintech company reproduced a production bug by writing a failing pytest that mimicked the customer's inputs. After the fix, the test joined the regression suite, ensuring the issue could never return silently.

```
from __future__ import annotations

from billing.reports import render_statement

def test_regression_missing_currency() -> None:
    """Lock in the fix for an incident triggered by malformed input."""
    payload = render_statement(currency="EUR", amount=0)
    assert "currency" in payload
```

Listing 7.5: Regression test distilled from an incident

7.6 Summary

Healthy suites blend fast unit tests with targeted integrations, emphasise testable code structure, and capture regressions the moment they appear in CI pipelines.

```
from __future__ import annotations

def summarize_flakes(failing: list[str]) -> str:
    """Produce a string for Slack alerts describing flaky tests."""
    if not failing:
        return "No flaky tests detected this week."
    return f"Flaky tests: {', '.join(sorted(failing))}"
```

Listing 7.6: Summarising flaky tests

Exercises

1. Rewrite Listing 7.4 into a testable design with injected dependencies.
2. Add a property-based test using `hypothesis` for a function with tricky input domains.
3. Introduce contract tests for an external API your service consumes.
4. Configure coverage reporting in CI. Identify untested modules.
5. Design a rollback drill: break a test intentionally, fix it, and measure the time it takes.
6. Extend Listing 7.2 so it emits Prometheus metrics for each test classification.
7. Pair-review a flaky test. Capture the timeline of failures, hypothesise the root cause, and document the outcome in the regression suite.

```
from __future__ import annotations

from collections.abc import Iterable
from uuid import uuid4

def sync_documents(paths: Iterable[str], client, mover) -> None:
    """Delegate I/O to collaborators so tests can substitute them."""
    for path in paths:
        if path.endswith(".json"):
            client.upload_file(path, "analytics-bucket", f"imports/{uuid4()}")
        mover.move(path, "/archive")
```

Listing 7.7: Improved design with injectable collaborators

Chapter 8

Error Handling and Logging

8.1 Chapter Overview

Error handling and observability are inseparable. Without clear error pathways and structured logs, diagnosing issues becomes guesswork.

```
from __future__ import annotations

def categorize_error(exc: Exception) -> str:
    """Return an alert bucket given an exception instance."""
    name = exc.__class__.__name__
    if "Timeout" in name:
        return "transient"
    if "Validation" in name:
        return "user"
    return "system"
```

Listing 8.1: Categorising errors for alerting

8.2 Domain-Specific Exceptions

Define exceptions that reflect your business language so logs and alerts remain meaningful. Listing 4.6 already illustrated raising domain errors from services.

```
from __future__ import annotations

class BillingError(RuntimeError):
    """Base error for billing workflows."""

class PaymentInstrumentDeclined(BillingError):
    """Raised when the processor rejects a card."""
```

Listing 8.2: Expressive domain-level exception hierarchy

8.3 Structured Logging

```

import logging
from logging.config import dictConfig

def configure_logging() -> None:
    dictConfig(
        {
            "version": 1,
            "formatters": {
                "json": {
                    "format": "%(asctime)s %(levelname)s %(name)s %(message)s",
                    "class": "pythonjsonlogger.jsonlogger.JsonFormatter",
                }
            },
            "handlers": {
                "stderr": {
                    "class": "logging.StreamHandler",
                    "stream": "ext://sys.stderr",
                    "formatter": "json",
                }
            },
            "root": {"level": "INFO", "handlers": ["stderr"]},
        }
    )

logger = logging.getLogger("billing.events")
logger = logging.LoggerAdapter(logger, extra={"service": "billing"})

def log_payment_event(invoice_id: str, amount: str) -> None:
    logger.info("payment_completed", extra={"invoice_id": invoice_id,
                                              "amount": amount})

```

Listing 8.3: Structured logging with contextual metadata

8.4 Bad vs Good Error Handling

```

def run_job() -> None:
    try:
        do_work()
        persist_results()
    except Exception:
        return None

```

Listing 8.4: Anti-pattern: swallowing every exception

```

from __future__ import annotations

```

```

import logging

def run_job(logger: logging.Logger) -> None:
    try:
        do_work()
        persist_results()
    except NetworkError:
        logger.warning("Transient network issue; job will retry")
        raise
    except PersistenceError:
        logger.error("Failed to persist results", exc_info=True)
        raise

```

Listing 8.5: Improved version with explicit logging

8.5 Scenario: Incident Response

During an outage, the on-call engineer used structured logs to correlate failed HTTP requests with a downstream dependency. Because each log entry included contextual fields such as user ID and feature flag state, the team restored service within minutes.

```

from __future__ import annotations

from datetime import datetime

def build_incident_timeline(entries: list[dict[str, str]]) -> list[str]:
    """Return a human-readable sequence of incident events."""
    sorted_entries = sorted(entries, key=lambda entry: entry["timestamp"])
    return [
        f"{datetime.fromisoformat(entry['timestamp'])}: {entry['message']}"
        for entry in sorted_entries
    ]

```

Listing 8.6: Extracting timeline data from logs

8.6 Summary

Name your exceptions after real business events, log them with structure and context, and practice incident drills so the information proves useful under pressure.

```

from __future__ import annotations

def summarize_alerts(alerts: dict[str, int]) -> str:
    """Condense alert tallies into a planning sentence."""
    worst = max(alerts, key=alerts.get)
    return f"{{sum(alerts.values())}} alerts fired; {worst} noisy category."

```

Listing 8.7: Summarising alert volume

Exercises

1. Replace bare `except` blocks in your codebase with targeted exceptions.
2. Configure structured logging locally and confirm that observability dashboards parse the fields.
3. Write a chaos test that injects failures into a dependency and verify the logs are actionable.
4. Create a "runbook" log template that every subsystem uses for critical events.
5. Review a recent incident report and map each detection or mitigation step to logging improvements.
6. Extend Listing 8.6 so it groups events by feature flag and produces a Markdown report suitable for postmortems.

Chapter 9

Configuration and Secrets Management

9.1 Chapter Overview

Mismanaged configuration causes painful outages and leaks. Treat configuration as data, with types, validation, and lifecycle controls.

```
from __future__ import annotations

from dataclasses import dataclass
from datetime import datetime

@dataclass(slots=True)
class ConfigVersion:
    """Capture when and why configuration changed."""

    checksum: str
    applied_at: datetime
    change_reason: str
```

Listing 9.1: Tracking configuration versions

9.2 Typed Configuration Loader

```
from __future__ import annotations

from pathlib import Path

from pydantic import BaseSettings, Field


class Settings(BaseSettings):
    database_url: str = Field(..., env="DATABASE_URL")
    log_level: str = Field("INFO", env="LOG_LEVEL")
    feature_flag_path: Path = Field(Path("/etc/app/features.json"), env="FEATURE_FLAG_PATH")
```

```

class Config:
    env_file = ".env"

def load_settings() -> Settings:
    return Settings()

```

Listing 9.2: Single entry point for configuration

9.3 Secrets Discipline

Never commit credentials. Use environment variables, secret stores, or orchestrator injections. Run secret scanners like `trufflehog` in CI to catch mistakes early.

```

from __future__ import annotations

import os

def ensure_secret(name: str) -> str:
    """Raise an error if a required secret is absent."""
    if value := os.getenv(name):
        return value
    raise RuntimeError(f"Missing secret: {name}")

```

Listing 9.3: Validating secrets at startup

9.4 Scenario: Leaked Sandbox Key

A developer committed sandbox credentials that looked harmless. A malicious user discovered the repo and pivoted into production through a little-known trust relationship. After the incident, the company adopted automated scanning, rotated all credentials, and required code reviews for configuration files.

```

from __future__ import annotations

import subprocess

def scan_history() -> None:
    """Run a lightweight history scan before pushing."""
    subprocess.run(["git", "secrets", "--scan-history"], check=True)

```

Listing 9.4: Scanning history for secrets

9.5 Summary

Centralise configuration behind typed loaders, enforce secret hygiene, and document rotation and validation processes so deployments remain predictable.

```
from __future__ import annotations

def summarize_configs(configs: dict[str, bool]) -> str:
    """Report which services have adopted typed loaders."""
    adopted = [name for name, typed in configs.items() if typed]
    return f"{len(adopted)} services typed config; roll out remaining soon\n    ."
```

Listing 9.5: Summarising config coverage

Exercises

1. Refactor configuration access to funnel through a typed loader similar to Listing 9.2.
2. Add secret scanning to your CI pipeline and test it by committing a fake key.
3. Document the rotation process for every secret your service relies on.
4. Implement runtime validation that refuses to start when required environment variables are missing.
5. Design a feature flag strategy (file-based, database, or SaaS) and evaluate trade-offs.
6. Extend Listing 9.3 so it fetches secrets from your team's vault provider with caching and metrics hooks.

Chapter 10

Performance and Scalability

10.1 Chapter Overview

Optimise only after measuring. This chapter describes profiling, data structure choices, and concurrency strategies.

```
from __future__ import annotations

def regression_ratio(old_ms: float, new_ms: float) -> float:
    """Compute slowdown factor for benchmarking dashboards."""
    return round(new_ms / old_ms, 2)
```

Listing 10.1: Capturing performance regressions

10.2 Profiling Before Optimising

```
import cProfile
import pstats

from billing.reports import build_monthly_report

def profile_report() -> None:
    with cProfile.Profile() as profiler:
        build_monthly_report(project_id="alpha")
    stats = pstats.Stats(profiler)
    stats.sort_stats("cumulative").print_stats(15)
```

Listing 10.2: CProfile harness for a reporting job

10.3 Data Structures

Choose data structures that align with access patterns: `deque` for queues, `set` for membership tests, `heapq` for priority queues, and `itertools` for streaming. Avoid expensive operations such as repeated string concatenation inside loops.

```
from __future__ import annotations

from collections import deque

def process_queue(items: list[int]) -> deque[int]:
    """Convert lists into deques when pop-left operations dominate."""
    queue: deque[int] = deque(items)
    while queue and queue[0] < 0:
        queue.popleft()
    return queue
```

Listing 10.3: Selecting the right container

10.4 Concurrency Choices

Use `asyncio` for I/O-bound services, threads for blocking I/O when async refactors are impractical, and processes or native extensions for CPU-bound workloads. Protect shared state with locks or adopt immutable structures.

```
from __future__ import annotations

import asyncio

async def fetch_record(record_id: str) -> str:
    """Pretend to call an external service."""
    await asyncio.sleep(0.1)
    return record_id

async def fetch_many(ids: list[str]) -> list[str]:
    """Drive multiple calls concurrently for throughput."""
    return await asyncio.gather(*(fetch_record(item) for item in ids))
```

Listing 10.4: Running I/O-bound tasks concurrently

10.5 Scenario: Scaling a Report Generator

A consulting firm profiled its overnight PDF generator and discovered 80% of time spent in template rendering. By caching parsed templates and using `ThreadPoolExecutor` for I/O-bound API calls, they cut runtime from hours to minutes.

```
from __future__ import annotations

from functools import lru_cache

@lru_cache(maxsize=32)
def compile_template(name: str) -> str:
    """Simulate the expensive part of rendering."""
```

```
    return name.upper()
```

Listing 10.5: Caching template compilation

10.6 Summary

Let measurements guide every optimisation, choose data structures intentionally, and pick concurrency models that match the workload's constraints.

```
from __future__ import annotations

def summarize_speedups(results: dict[str, float]) -> str:
    """Generate a concise update for stakeholders."""
    fastest = min(results, key=results.get)
    return f"{fastest} achieved best runtime at {results[fastest]:.2f}s."
```

Listing 10.6: Summarising optimisation impact

Exercises

1. Profile a slow task using `cProfile`. Optimise the top offending function and measure the difference.
2. Rewrite a loop that concatenates strings with a more efficient approach using `"".join()`.
3. Implement both threaded and async versions of a simple downloader. Compare code complexity and throughput.
4. Build a benchmark harness that records baseline metrics for a key workflow.
5. Identify a cacheable computation and design an invalidation strategy.
6. Extend Listing 10.4 so it enforces a concurrency limit and records latency histograms.

Chapter 11

Security for Python Developers

11.1 Chapter Overview

Security is everyone's responsibility. This chapter addresses dependency hygiene, input validation, and safe filesystem access.

```
from __future__ import annotations

def compute_vulnerability_rate(vulnerable: int, total: int) -> float:
    """Return proportion of dependencies with known issues."""
    return round(vulnerable / total, 2)
```

Listing 11.1: Capturing CVE exposure

11.2 Dependency Hygiene

Run pip-audit or Safety regularly. Automate updates with Dependabot or Renovate so vulnerabilities are patched quickly.

```
from __future__ import annotations

import subprocess

def run_pip_audit() -> None:
    """Fail the build when vulnerabilities exist."""
    subprocess.run(["pip-audit", "--strict"], check=True)
```

Listing 11.2: Invoking pip-audit from Python

11.3 Input Validation

Use Pydantic models, argparse, or custom validators to treat all external input as untrusted. Explicitly constrain formats and ranges.

```
from __future__ import annotations
```

```
from pydantic import BaseModel, Field

class PaymentRequest(BaseModel):
    """Trusted structure for inbound payloads."""

    account_id: str = Field(min_length=8, max_length=32)
    amount_cents: int = Field(gt=0, lt=100_000_00)
```

Listing 11.3: Validating API payloads

11.4 Safe File Handling

```
from __future__ import annotations

from pathlib import Path

def read_report(report_name: str, base_dir: Path) -> str:
    reports_dir = base_dir / "reports"
    reports_dir.mkdir(exist_ok=True)
    candidate = (reports_dir / report_name).resolve()
    if reports_dir.resolve() not in candidate.parents:
        raise PermissionError("Illegal path traversal attempt detected")
    return candidate.read_text(encoding="utf-8")
```

Listing 11.4: Safe path resolution guards against traversal

11.5 Scenario: Dependency Supply-Chain Attack

An open-source package added a malicious post-install hook. Because the team pinned versions and used hash-checking installers, the compromised release never reached production. They then set up signed releases and mirrored dependencies internally.

```
from __future__ import annotations

import hashlib
from pathlib import Path

def verify_package(path: Path, expected_hash: str) -> None:
    """Raise when downloaded artifacts do not match expectation."""
    digest = hashlib.sha256(path.read_bytes()).hexdigest()
    if digest != expected_hash:
        raise RuntimeError("Package hash mismatch detected")
```

Listing 11.5: Validating package hashes

11.6 Summary

Keep dependencies patched, treat every external input as hostile, and harden filesystem access to avoid trivial escalation paths.

```
from __future__ import annotations

def summarize_findings(findings: list[str]) -> str:
    """Condense security findings for exec summaries."""
    if not findings:
        return "No blocking security findings."
    critical = [item for item in findings if "critical" in item.lower()]
    return f"{len(findings)} findings ({len(critical)}) critical remain."
```

Listing 11.6: Summarising security posture

Exercises

1. Run `pip-audit` on your project and triage any findings.
2. Add validation for a CLI command that currently trusts user input.
3. Design a security review checklist for third-party libraries.
4. Implement hash checking for dependency installation using `pip -require-hashes`.
5. Build a small demo exploiting a path traversal bug, then patch it as in Listing 11.4.
6. Extend Listing 11.5 so it also validates a digital signature before trusting a binary package.

Chapter 12

Packaging, Distribution, and Versioning

12.1 Chapter Overview

Packaging turns code into artifacts that other teams can trust. Follow modern packaging standards and disciplined versioning.

```
from __future__ import annotations

from dataclasses import dataclass

@dataclass(slots=True)
class Release:
    """Track release metadata for internal dashboards."""

    version: str
    commit: str
    released_by: str
```

Listing 12.1: Recording release metadata

12.2 Build Real Packages

Provide `__init__.py` files, export stable APIs, and publish wheels to internal or public registries. Modern workflows revolve around `pyproject.toml`.

```
from __future__ import annotations

import subprocess

def build_package() -> None:
    """Invoke the build backend and emit sdist + wheel."""
    subprocess.run(["python", "-m", "build"], check=True)
```

Listing 12.2: Building distributions via Python

12.3 Entry Points

```
[project.scripts]
bill = "billing.cli:main"
```

Listing 12.3: Console script entry point

12.4 Versioning Discipline

Adopt semantic versioning: MAJOR for breaking changes, MINOR for new features, PATCH for bug fixes. Tag releases in Git and maintain a human-readable changelog.

```
from __future__ import annotations

def bump(version: str, level: str) -> str:
    """Return a new semver string at the requested level."""
    major, minor, patch = map(int, version.split('.'))
    match level:
        case "major":
            major += 1
            minor = 0
            patch = 0
        case "minor":
            minor += 1
            patch = 0
        case "patch":
            patch += 1
        case _:
            raise ValueError("Unknown level")
    return f"{major}.{minor}.{patch}"
```

Listing 12.4: Bumping semantic versions programmatically

12.5 Scenario: Coordinated Release Train

A platform team shipped a library consumed by ten services. They instituted a release train where every Wednesday a new minor version shipped with release notes and migration guides. Incidents dropped because downstream teams could plan upgrades.

```
from __future__ import annotations

from datetime import date, timedelta

def release_train(start: date, *, cadence_days: int = 7) -> list[date]:
    """Return upcoming release windows."""
    return [start + timedelta(days=cadence_days * offset) for offset in
           range(4)]
```

Listing 12.5: Generating release train schedules

12.6 Summary

Ship libraries with modern metadata, expose clear entry points, and manage semantic versions plus changelogs so downstream consumers can upgrade confidently.

```
from __future__ import annotations

def summarize_consumers(consumers: list[str]) -> str:
    """Return a summary of services pinned to the latest release."""
    return f"{len(consumers)} downstream services upgraded this sprint."
```

Listing 12.6: Summarising adoption

Exercises

1. Package a small module into a wheel and install it in a clean environment.
2. Create a changelog entry for a hypothetical breaking change.
3. Add release automation to CI (e.g., publish to TestPyPI on tag).
4. Define compatibility guarantees for your public API and document them.
5. Evaluate whether your project should use namespace packages or a monorepo layout.
6. Extend Listing 12.4 so it updates `pyproject.toml` and commits the change automatically.

```
from __future__ import annotations

def main() -> None:
    """Entrypoint invoked by console scripts."""
    print("Billing CLI ready")
```

Listing 12.7: Implementing the console script target

Chapter 13

Collaboration, Git, and CI/CD

13.1 Chapter Overview

People practices sustain technical excellence. This chapter covers source control habits, code reviews, and automation.

```
from __future__ import annotations

def review_latency(hours_waited: list[int]) -> float:
    """Return the median wait time for reviews."""
    sorted_hours = sorted(hours_waited)
    mid = len(sorted_hours) // 2
    return float(sorted_hours[mid])
```

Listing 13.1: Tracking collaboration metrics

13.2 Git Hygiene

Commit early with focused changes. Write imperative commit messages. Rebase feature branches onto main before opening pull requests.

```
from __future__ import annotations

def validate_commit_message(message: str) -> None:
    """Ensure commits follow imperative style."""
    if not message or message[0].islower():
        raise SystemExit("Commit message must start with an imperative verb.")
```

Listing 13.2: Guarding commit messages via hook

13.3 Code Review Culture

Treat reviews as collaborative design conversations. Authors provide context and testing evidence; reviewers prioritise correctness and maintainability.

```
from __future__ import annotations

def build_pr_template(tests: list[str], risks: list[str]) -> str:
    """Produce a review-ready description."""
    tests_block = "\n".join(f"- {test}" for test in tests)
    risks_block = "\n".join(f"- {item}" for item in risks)
    return f"## Tests\n{tests_block}\n\n## Risks\n{risks_block}"
```

Listing 13.3: Summarising review context automatically

13.4 Automation and Hooks

```
repos:
  - repo: https://github.com/charliermarsh/ruff-pre-commit
    rev: v0.6.8
    hooks:
      - id: ruff
      - id: ruff-format
  - repo: https://github.com/psf/black
    rev: 24.8.0
    hooks:
      - id: black
  - repo: https://github.com/pre-commit/mirrors-mypy
    rev: v1.11.0
    hooks:
      - id: mypy
```

Listing 13.4: Pre-commit configuration for consistent tooling

13.5 Scenario: CI as Gatekeeper

A data team configured CI to run linters, type checkers, tests, and deployment previews. A regression slipped through local testing but failed in CI because the environment built from scratch. The automated gate prevented a broken migration from reaching production.

```
from __future__ import annotations

def ensure_checks(checks: dict[str, bool]) -> None:
    """Guard CI merges until each check reports success."""
    missing = [name for name, passed in checks.items() if not passed]
    if missing:
        raise RuntimeError(f"CI gatekeeper blocked by {missing}")
```

Listing 13.5: Failing fast when CI requirements are missing

13.6 Summary

Healthy collaboration relies on disciplined git hygiene, respectful reviews, and CI pipelines that run the same automated checks for everyone.

```
from __future__ import annotations

def summarize_collab(metrics: dict[str, float]) -> str:
    """Return a narrative summary for the weekly sync."""
    return (
        f"Median review wait {metrics['review_latency']}h; "
        f"{metrics['ci_pass_rate']} * 100:.0f}% CI pass rate."
    )
```

Listing 13.6: Summarising collaboration signals

Exercises

1. Review your last five commits. Were they cohesive? If not, how would you split them?
2. Establish a rotating reviewer schedule to spread knowledge.
3. Add `pre-commit` to your repository and enforce it in CI.
4. Create a checklist for pull request descriptions (context, tests, screenshots).
5. Simulate a CI outage and document manual fallback steps.
6. Extend Listing 13.5 so it posts a Slack alert whenever a gatekeeper blocks a merge.

```
from __future__ import annotations

import subprocess

def run_hooks() -> None:
    """Re-run hooks locally so CI never surprises engineers."""
    subprocess.run(["pre-commit", "run", "--all-files"], check=True)
```

Listing 13.7: Python helper to run pre-commit

Chapter 14

Case Studies and Patterns in Real Projects

14.1 Chapter Overview

Abstract guidance becomes tangible through stories. This chapter captures recurring patterns from real teams.

```
from __future__ import annotations

def case_study_summary(name: str, impact: str) -> str:
    """Return a short descriptor for dashboards."""
    return f"{name}: {impact}"
```

Listing 14.1: Representing case study metadata

14.2 Billing Platform Modernisation

A SaaS billing team inherited a cron-based script that generated invoices. They introduced source layout discipline (Listing 4.2), wrapped calculations with typed modules (Listing 2.6), and added pytest coverage for every pricing rule. Structured logging (Listing 8.3) helped support staff diagnose customer complaints. Deployment frequency increased because engineers trusted their safety nets.

```
from __future__ import annotations

from dataclasses import dataclass

@dataclass(slots=True)
class InvoiceProjection:
    """Aggregate data used during the refactor."""

    customer_id: str
    subtotal: float
    taxes: float
```

```
def total_due(projection: InvoiceProjection) -> float:
    """Return the amount customers see post-refactor."""
    return projection.subtotal + projection.taxes
```

Listing 14.2: Capturing invoice recalculation state

14.3 Data Pipeline Hardening

An analytics squad maintained a nightly ingestion job similar to Listing 7.4. After a costly outage, they refactored into injectable components, added property-based tests to validate CSV parsing, and enforced configuration loading via `Settings`. CI now spins up ephemeral storage backends for integration tests, and performance monitoring from Chapter 10 alerts engineers when runtimes drift.

```
from __future__ import annotations

from collections.abc import Iterable

def transform_rows(rows: Iterable[str], *, delimiter: str = ",") -> list[list[str]]:
    """Ensure pipeline logic works with dependency injection."""
    return [row.split(delimiter) for row in rows]
```

Listing 14.3: Injectable pipeline components

14.4 Future Work

% TODO: Add a machine-learning-focused case study covering experiment tracking, reproducible data, and model deployment.

```
from __future__ import annotations

def placeholder_case(title: str) -> dict[str, str]:
    """Reserve space for future patterns."""
    return {"title": title, "status": "draft"}
```

Listing 14.4: Scaffolding future case studies

14.5 Summary

Case studies show how layered practices—structure, typing, testing, observability—combine to deliver predictable outcomes, and they highlight where future research is needed.

```
from __future__ import annotations

def summarize_lessons(lessons: list[str]) -> str:
    """Produce a concise statement for leadership debriefs."""
    return "; ".join(lessons[:3])
```

Listing 14.5: Summarising lessons learned

Exercises

1. Interview another team about their most successful refactor. Map their steps to the practices in earlier chapters.
2. Create a mini case study for your project, documenting before/after metrics.
3. Design a "playbook" template for capturing future case studies.
4. Identify a legacy system in your organisation and propose the first three incremental steps toward modernisation.
5. Reflect on a failed initiative. Which missing practices contributed to the outcome?
6. Implement code based on Listing 14.7 that validates CSV headers before transformation to avoid future ingestion incidents.

```
from __future__ import annotations

from dataclasses import dataclass


@dataclass(slots=True)
class InvoiceProjection:
    """Aggregate data used during the refactor."""

    customer_id: str
    subtotal: float
    taxes: float

    def total_due(projection: InvoiceProjection) -> float:
        """Return the amount customers see post-refactor."""
        return projection.subtotal + projection.taxes
```

Listing 14.6: Capturing invoice recalculation state

```
from __future__ import annotations

from collections.abc import Iterable


def transform_rows(rows: Iterable[str], *, delimiter: str = ",") -> list[list[str]]:
    """Ensure pipeline logic works with dependency injection."""
    return [row.split(delimiter) for row in rows]
```

Listing 14.7: Injectable pipeline components

Appendix A

Checklists, Templates, and Further Reading

A.1 Chapter Overview

The appendix collects reusable assets that keep teams aligned.

```
from __future__ import annotations

def checklist_owner(name: str) -> str:
    """Map checklist names to owners."""
    return f"{name} owner: platform team"
```

Listing A.1: Representing checklist ownership

A.2 Launch Checklist

Before shipping, confirm that formatting, linting, typing, and tests pass locally and in CI. Review documentation updates, rotate credentials if needed, and run deployment rehearsals.

```
from __future__ import annotations

def launch_ready(checks: dict[str, bool]) -> bool:
    """Return True when every checklist item passes."""
    return all(checks.values())
```

Listing A.2: Evaluating launch readiness programmatically

A.3 Template Repository

Maintain a starter repo containing `pyproject.toml`, `noxfile.py`, CI workflows, and documentation scaffolding. New services can fork it to inherit best practices.

```
from __future__ import annotations

import shutil
```

```
from pathlib import Path

def scaffold_from_template(template: Path, destination: Path) -> None:
    """Copy the template repo into a new project directory."""
    shutil.copytree(template, destination, dirs_exist_ok=True)
```

Listing A.3: Scaffolding a new repository from a template

A.4 Further Reading

Recommended resources include Brett Slatkin’s *Effective Python*, Hynek Schlawack’s *Solid Python*, the official *Python Packaging User Guide*, and the *Twelve Factor App*. Follow tool maintainers (Black, Ruff, Pytest, Mypy) to stay current.

```
from __future__ import annotations

def reading_plan(topics: list[str]) -> list[str]:
    """Pair topics with recommended books."""
    return [f"{topic}: Effective Python (2nd Ed.)" for topic in topics]
```

Listing A.4: Generating reading lists

A.5 Summary

Standardised checklists, templates, and curated resources ensure every new initiative starts with the same proven guardrails.

```
from __future__ import annotations

def summarize_assets(counts: dict[str, int]) -> str:
    """Provide a short statement listing asset counts."""
    return ", ".join(f"{kind}: {count}" for kind, count in counts.items())
```

Listing A.5: Summarising appendix assets

Exercises

1. Build a personalised launch checklist and compare it with this appendix.
2. Create a template repository for your organisation and solicit feedback.
3. Curate a reading list tailored to your team’s domain (web, data, ML).
4. Run a brown-bag session where each engineer shares one tool configuration tip.
5. Evaluate whether your onboarding guide references every relevant checklist; update as needed.
6. Extend Listing A.2 so it prints which checklist item failed when the launch is blocked.